

SML 201 – Week 2

John D. Storey

Spring 2016

Getting Started in R

Summary from Week 1

Last week we learned about R:

- calculations
- getting help
- atomic classes
- assigning values to variables
- factors
- vectors, matrices, lists
- some basic functions

Missing Values

In data analysis and model fitting, we often have missing values. NA represents missing values and NaN means “not a number”, which is a special type of missing value.

```
> m <- matrix(nrow=3, ncol=3)
> m
      [,1] [,2] [,3]
[1,]    NA    NA    NA
[2,]    NA    NA    NA
[3,]    NA    NA    NA
> 0/1
[1] 0
> 1/0
[1] Inf
> 0/0
[1] NaN
```

NULL

NULL is a special type of reserved value in R.

```
> x <- vector(mode="list", length=3)
> x
[[1]]
NULL

[[2]]
NULL

[[3]]
NULL
```

Coercion

We saw earlier that when we mixed classes in a vector they were all coerced to be of type character:

```
> c("a", 1:3, TRUE, FALSE)
[1] "a"      "1"      "2"      "3"      "TRUE"   "FALSE"
```

You can directly apply coercion with functions as `.numeric()`, `.character()`, `.logical()`, etc.

This doesn't always work out well:

```
> x <- 1:3
> as.character(x)
[1] "1" "2" "3"
>
> y <- c("a", "b", "c")
> as.numeric(y)
Warning: NAs introduced by coercion
[1] NA NA NA
```

Lists (review)

Lists allow you to hold different classes of objects in one variable.

```
> x <- list(1:3, "a", c(TRUE, FALSE))
> x
[[1]]
[1] 1 2 3

[[2]]
[1] "a"

[[3]]
[1] TRUE FALSE
>
> # access any element of the list
> x[[2]]
[1] "a"
> x[[3]][2]
[1] FALSE
```

Lists with Names (review)

The elements of a list can be given names.

```
> x <- list(counting=1:3, char="a", logic=c(TRUE, FALSE))
> x
$counting
[1] 1 2 3

$char
[1] "a"

$logic
[1] TRUE FALSE
>
> # access any element of the list
> x$char
[1] "a"
> x$logic[2]
[1] FALSE
```


Data Frames

The data frame is one of the most important objects in R. Data sets very often come in tabular form of mixed classes, and data frames are constructed exactly for this.

Data frames are lists where each element has the same length.

Data Frames

```
> df <- data.frame(counting=1:3, char=c("a", "b", "c"),  
+                  logic=c(TRUE, FALSE, TRUE))  
> df  
  counting char logic  
1         1    a  TRUE  
2         2    b FALSE  
3         3    c  TRUE  
>  
> nrow(df)  
[1] 3  
> ncol(df)  
[1] 3
```

Data Frames

```
> dim(df)
[1] 3 3
>
> names(df)
[1] "counting" "char"      "logic"
>
> attributes(df)
$names
[1] "counting" "char"      "logic"

$row.names
[1] 1 2 3

$class
[1] "data.frame"
```

Attributes

Attributes give information (or meta-data) about R objects. The previous slide shows `attributes(df)`, the attributes of the data frame `df`.

```
> x <- 1:3
> attributes(x) # no attributes for a standard vector
NULL
>
> m <- matrix(1:6, nrow=2, ncol=3)
> attributes(m)
$dim
[1] 2 3
```

Attributes (cont'd)

```
> paint <- factor(c("red", "white", "blue", "blue", "red",  
+                  "red"))  
> attributes(paint)  
$levels  
[1] "blue"  "red"   "white"  
  
$class  
[1] "factor"
```

Names

Names can be assigned to columns and rows of vectors, matrices, and data frames. This makes your code easier to write and read.

```
> names(x) <- c("Princeton", "Rutgers", "Penn")
> x
Princeton    Rutgers      Penn
         1         2         3
>
> colnames(m) <- c("NJ", "NY", "PA")
> rownames(m) <- c("East", "West")
> m
      NJ NY PA
East  1  3  5
West  2  4  6
> colnames(m)
[1] "NJ" "NY" "PA"
```

Accessing Names

Displaying or assigning names to these three types of objects does not have consistent syntax.

Object	Column Names	Row Names
vector	<code>names ()</code>	N/A
data frame	<code>names ()</code>	<code>row.names ()</code>
data frame	<code>colnames ()</code>	<code>rownames ()</code>
matrix	<code>colnames ()</code>	<code>rownames ()</code>

Reproducibility

Definition and Motivation

- Reproducibility involves *being able to recalculate the exact numbers in a data analysis using the code and raw data provided by the analyst.*
- Reproducibility is often difficult to achieve and has slowed down the discovery of important data analytic errors.
- Reproducibility should not be confused with “correctness” of a data analysis. A data analysis can be fully reproducible and recreate all numbers in an analysis and still be misleading or incorrect.

Taken from *Elements of Data Analytic Style*

Reproducible vs. Replicable

Reproducible research is often used these days to indicate the ability to recalculate the exact numbers in a data analysis

Replicable research results often refers to the ability to independently carry out a study (thereby collecting new data) and coming to equivalent conclusions as the original study

These two terms are often confused, so it is important to clearly state the definition

Steps to a Reproducible Analysis

1. Use a data analysis script – e.g., R Markdown (discussed next section!) or iPython Notebooks
2. Record versions of software and parameters – e.g., use `sessionInfo()` in R as in `project_1.Rmd`
3. Organize your data analysis
4. Use version control – e.g., GitHub
5. Set a random number generator seed – e.g., use `set.seed()` in R
6. Have someone else run your analysis

Organizing Your Data Analysis

- Data
 - raw data
 - processed data (sometimes multiple stages for very large data sets)
- Figures
 - Exploratory figures
 - Final figures

Organizing Your Data Analysis (cont'd)

- R code
 - Raw or unused scripts
 - Data processing scripts
 - Analysis scripts
- Text
 - README files explaining what all the components are
 - Final data analysis products like presentations/writeups

Common Mistakes

- Failing to use a script for your analysis
- Not recording software and package version numbers or other settings used
- Not sharing your data and code
- Using reproducibility as a social weapon

R Markdown

R + Markdown + knitr

R Markdown was developed by the RStudio team to allow one to write reproducible research documents using Markdown and `knitr`. This is contained in the `rmarkdown` package, but can easily be carried out in RStudio.

Markdown was originally developed as a very simply text-to-html conversion tool. With Pandoc, Markdown is a very simply text-to-x conversion tool where x can be many different formats: html, LaTeX, PDF, Word, etc.

R Markdown Files

R Markdown documents begin with a metadata section, the YAML header, that can include information on the title, author, and date as well as options for customizing output.

```
---  
title: "SML 201 -- Project 1"  
author: "Your Name"  
date: February 8, 2016  
output: pdf_document  
---
```

Many options are available. See [**http://rmarkdown.rstudio.com**](http://rmarkdown.rstudio.com) for full documentation.

Markdown

Headers:

```
# Header 1
## Header 2
### Header 3
```

Emphasis:

```
*italic* **bold**
_italic_ __bold__
```

Tables:

First Header	Second Header
Content Cell	Content Cell
Content Cell	Content Cell

Markdown (cont'd)

Unordered list:

- Item 1
- Item 2
 - Item 2a
 - Item 2b

Ordered list:

1. Item 1
2. Item 2
3. Item 3
 - Item 3a
 - Item 3b

Markdown (cont'd)

Links:

```
http://example.com
```

```
[linked phrase](http://example.com)
```

Blockquotes:

```
Florence Nightingale once said:
```

```
> For the sick it is important  
> to have the best.
```

Markdown (cont'd)

Plain code blocks:

```
```\nThis text is displayed verbatim with no formatting.\n```
```

Inline Code:

```
We use the `print()` function to print the contents\nof a variable in R.
```

Additional documentation and examples can be found **here** and **here**.

# knitr

The `knitr` R package allows one to execute R code within a document, and to display the code itself and its output (if desired). This is particularly easy to do in the R Markdown setting. For example...

*Placing the following text in an R Markdown file*

The sum of 2 and 2 is ``r 2+2``.

*produces in the output file*

The sum of 2 and 2 is 4.

# knitr Chunks

Chunks of R code separated from the text. In R Markdown:

```
```${r}  
x <- 2  
x + 1  
print(x)  
```
```

Output in file:

```
> x <- 2
> x + 1
[1] 3
> print(x)
[1] 2
```

# Chunk Option: echo

In R Markdown:

```
````{r, echo=FALSE}  
x <- 2  
x + 1  
print(x)  
````
```

Output in file:

```
[1] 3
[1] 2
```



# Chunk Option: results

In R Markdown:

```
```{r, results="hide"}  
x <- 2  
x + 1  
print(x)  
```
```

Output in file:

```
> x <- 2
> x + 1
> print(x)
```

# Chunk Option: `include`

In R Markdown:

```
```{r, include=FALSE}  
x <- 2  
x + 1  
print(x)  
```
```

Output in file:

(nothing)

# Chunk Option: eval

In R Markdown:

```
```{r, eval=FALSE}  
x <- 2  
x + 1  
print(x)  
```
```

Output in file:

```
> x <- 2
> x + 1
> print(x)
```

# Chunk Names

Naming your chunks can be useful for identifying them in your file and during the execution, and also to denote dependencies among chunks.

```
```{r my_first_chunk}  
x <- 2  
x + 1  
print(x)  
```
```

# knitr Option: cache

Sometimes you don't want to run chunks over and over, especially for large calculations. You can “cache” them.

```
```{r chunk1, cache=TRUE, include=FALSE}  
x <- 2  
```
```

```
```{r chunk2, cache=TRUE, dependson='chunk1'}  
y <- 3  
z <- x + y  
```
```

This creates a directory called `cache` in your working directory that stores the objects created or modified in these chunks. When `chunk1` is modified, it is re-run. Since `chunk2` depends on `chunk1`, it will also be re-run.

# knitr Options: figures

You can add chunk options regarding the placement and size of figures.  
Examples include:

- `fig.width`
- `fig.height`
- `fig.align`

# Changing Default Chunk Settings

If you will be using the same options on most chunks, you can set default options for the entire document. Run something like this at the beginning of your document with your desired chunk options.

```
```{r my_opts, cache=FALSE, echo=FALSE}  
library(knitr)  
opts_chunk$set(fig.align="center", fig.height=4, fig.width=6)  
```
```

# Documentation and Examples

- <http://yihui.name/knitr/>
- [http://kbroman.org/knitr\\_knutshell/pages/Rmarkdown.html](http://kbroman.org/knitr_knutshell/pages/Rmarkdown.html)
- <https://github.com/SML201/lectures>



# Control Structures

# Rationale

- Control structures in R allow you to control the flow of execution of a series of R expressions
- They allow you to put some logic into your R code, rather than just always executing the same R code every time
- Control structures also allow you to respond to inputs or to features of the data and execute different R expressions accordingly

Paraphrased from *R Programming for Data Science*

# Common Control Structures

- `if` and `else`: testing a condition and acting on it
- `for`: execute a loop a fixed number of times
- `while`: execute a loop while a condition is true
- `repeat`: execute an infinite loop (must break out of it to stop)
- `break`: break the execution of a loop
- `next`: skip an iteration of a loop

From *R Programming for Data Science*

# Some Boolean Logic

R has built-in functions that produce TRUE or FALSE such as `is.vector` or `is.na`. You can also do the following:

- `x == y` : does x equal y?
- `x > y` : is x greater than y? (also `<` less than)
- `x >= y` : is x greater than or equal to y?
- `x && y` : are both x and y true?
- `x || y` : is either x or y true?
- `!is.vector(x)` : this is TRUE if x is not a vector

# if

Idea:

```
if(<condition>) {
 ## do something
}
Continue with rest of code
```

Example:

```
> x <- c(1,2,3)
> if(is.numeric(x)) {
+ x+2
+ }
[1] 3 4 5
```

# if-else

Idea:

```
if(<condition>) {
 ## do something
}
else {
 ## do something else
}
```

Example:

```
> x <- c("a", "b", "c")
> if(is.numeric(x)) {
+ print(x+2)
+ } else {
+ class(x)
+ }
[1] "character"
```

# for Loops

Example:

```
> for(i in 1:10) {
+ print(i)
+ }
[1] 1
[1] 2
[1] 3
[1] 4
[1] 5
[1] 6
[1] 7
[1] 8
[1] 9
[1] 10
```

# for Loops (cont'd)

## Examples:

```
> x <- c("a", "b", "c", "d")
>
> for(i in 1:4) {
+ print(x[i])
+ }
[1] "a"
[1] "b"
[1] "c"
[1] "d"
>
> for(i in seq_along(x)) {
+ print(x[i])
+ }
[1] "a"
[1] "b"
[1] "c"
[1] "d"
```



# Nested for Loops

Example:

```
> m <- matrix(1:6, nrow=2, ncol=3, byrow=TRUE)
>
> for(i in seq_len(nrow(m))) {
+ for(j in seq_len(ncol(m))) {
+ print(m[i,j])
+ }
+ }
[1] 1
[1] 2
[1] 3
[1] 4
[1] 5
[1] 6
```

# while

Example:

```
> x <- 1:10
> idx <- 1
>
> while(x[idx] < 4) {
+ print(x[idx])
+ idx <- idx + 1
+ }
[1] 1
[1] 2
[1] 3
>
> idx
[1] 4
```

Repeats the loop until while the condition is TRUE.

# repeat

Example:

```
> x <- 1:10
> idx <- 1
>
> repeat {
+ print(x[idx])
+ idx <- idx + 1
+ if(idx >= 4) {
+ break
+ }
+ }
[1] 1
[1] 2
[1] 3
>
> idx
[1] 4
```

Repeats the loop until break is executed.

# break and next

`break` ends the loop. `next` skips the rest of the current loop iteration.

Example:

```
> x <- 1:1000
> for(idx in 1:1000) {
+ # %% calculates division remainder
+ if((x[idx] %% 2) > 0) {
+ next
+ } else if(x[idx] > 10) { # an else-if!!
+ break
+ } else {
+ print(x[idx])
+ }
+ }
[1] 2
[1] 4
[1] 6
[1] 8
[1] 10
```

# Vectorized Operations

# Calculations on Vectors

R is usually smart about doing calculations with vectors. Examples:

```
>
> x <- 1:3
> y <- 4:6
>
> 2*x # same as c(2*x[1], 2*x[2], 2*x[3])
[1] 2 4 6
> x + 1 # same as c(x[1]+1, x[2]+1, x[3]+1)
[1] 2 3 4
> x + y # same as c(x[1]+y[1], x[2]+y[2], x[3]+y[3])
[1] 5 7 9
> x*y # same as c(x[1]*y[1], x[2]*y[2], x[3]*y[3])
[1] 4 10 18
```

# A Caveat

If two vectors are of different lengths, R tries to find a solution for you (and doesn't always tell you).

```
> x <- 1:5
> y <- 1:2
> x+y
Warning in x + y: longer object length is not a multiple of shorter object
length
[1] 2 4 4 6 6
```

What happened here?

# Vectorized Matrix Operations

Operations on matrices are also vectorized. Example:

```
> x <- matrix(1:4, nrow=2, ncol=2, byrow=TRUE)
> y <- matrix(1:4, nrow=2, ncol=2)
>
> x+y
 [,1] [,2]
[1,] 2 5
[2,] 5 8
>
> x*y
 [,1] [,2]
[1,] 1 6
[2,] 6 16
```



# Mixing Vectors and Matrices

What happens when we do calculations involving a vector and a matrix?

Example:

```
> x <- matrix(1:6, nrow=2, ncol=3, byrow=TRUE)
> z <- 1:2
>
> x + z
 [,1] [,2] [,3]
[1,] 2 3 4
[2,] 6 7 8
>
> x * z
 [,1] [,2] [,3]
[1,] 1 2 3
[2,] 8 10 12
```

# Mixing Vectors and Matrices

Another example:

```
> x <- matrix(1:6, nrow=2, ncol=3, byrow=TRUE)
> z <- 1:3
>
> x + z
 [,1] [,2] [,3]
[1,] 2 5 5
[2,] 6 6 9
>
> x * z
 [,1] [,2] [,3]
[1,] 1 6 6
[2,] 8 5 18
```

What happened this time?

# Vectorized Boolean Logic

We saw `&&` and `||` applied to pairs of logical values. We can also vectorize these operations.

```
> a <- c(TRUE, TRUE, FALSE)
> b <- c(FALSE, TRUE, FALSE)
>
> a | b
[1] TRUE TRUE FALSE
> a & b
[1] FALSE TRUE FALSE
```

# Functions

# Rationale

- Writing functions is a core activity of an R programmer. It represents the key step of the transition from a mere user to a developer who creates new functionality for R.
- Functions are often used to encapsulate a sequence of expressions that need to be executed numerous times, perhaps under slightly different conditions.
- Functions are also often written when code must be shared with others or the public.

From *R Programming for Data Science*

# Defining a New Function

- Functions are defined using the `function ( )` directive
- They are stored as variables, so they can be passed to other functions and assigned to new variables
- Arguments and a final return object are defined

# Example 1

```
> my_square <- function(x) {
+ x*x # can also do return(x*x)
+ }
>
> my_square(x=2)
[1] 4
>
> my_fun2 <- my_square
> my_fun2(x=3)
[1] 9
```

## Example 2

```
> my_square_ext <- function(x) {
+ y <- x*x
+ return(list(x_original=x, x_squared=y))
+ }
>
> my_square_ext(2)
$x_original
[1] 2

$x_squared
[1] 4
>
> z <- my_square_ext(2)
```



# Example 3

```
> my_power <- function(x, e, say_hello) {
+ if(say_hello) {
+ cat("Hello World!")
+ }
+ x^e
+ }
>
> my_power(x=2, e=3, say_hello=TRUE)
Hello World!
[1] 8
>
> z <- my_power(x=2, e=3, say_hello=TRUE)
Hello World!
> z
[1] 8
```

# Default Function Argument Values

Some functions have default values for their arguments:

```
> str(matrix)
function (data = NA, nrow = 1, ncol = 1, byrow = FALSE, dimnames = NULL)
```

You can define a function with default values by the following:

```
f <- function(x, y=2) {
 x + y
}
```

If the user types `f ( x=1 )` then it defaults to `y=2`, but if the user types `f ( x=1 , y=3 )`, then it executes with these assignments.

# The Ellipsis Argument

You will encounter functions that include as a possible argument the ellipsis: . . .

This basically holds arguments that can be passed to functions called within a function. Example:

```
> double_log <- function(x, ...) {
+ log(2*x, ...)
+ }
>
> double_log(1, base=2)
[1] 1
> double_log(1, base=10)
[1] 0.30103
```

# Argument Matching

*R Programming for Data Science* spends several pages discussing how R deals with function calls when the arguments are not defined explicitly. For example:

```
x <- matrix(1:6, nrow=2, ncol=3, byrow=TRUE) # versus
x <- matrix(1:6, 2, 3, TRUE)
```

I strongly recommend that you define arguments explicitly. For example, I can never remember which comes first in `matrix()`, `nrow` or `ncol`.

# Organizing Your Code

# Suggestions

Rstudio conveniently tries to automatically format your R code. We suggest the following in general.

1. No more than 80 characters per line (or fewer depending on how R Markdown compiles):

```
really_long_line <- my_function(x=20, y=30, z=TRUE,
 a="Joe", b=3.8)
```

2. Indent 2 or more characters for nested commands:

```
for(i in 1:10) {
 if(i > 4) {
 print(i)
 }
}
```

# Suggestions (cont'd)

## 3. Generously comment your code.

```
a for-loop that prints the index
whenever it is greater than 4
for(i in 1:10) {
 if(i > 4) {
 print(i)
 }
}
a good way to get partial credit
if something goes wrong :-)
```

## 4. Do not hesitate to write functions to organize tasks. These help to break up your code into more understandable pieces, and functions can often be used several times.

# Where to Put Files

See *Elements of Data Analytic Style*, Chapter 12 (“Reproducibility”) for suggestions on how to organize your files.

In this course, we will keep this relatively simple. We will try to provide you with some organization when distributing the projects.



# Packages

# Rationale

“In R, the fundamental unit of shareable code is the package. A package bundles together code, data, documentation, and tests, and is easy to share with others. As of January 2015, there were over 6,000 packages available on the **Comprehensive R Archive Network**, or **CRAN**, the public clearing house for R packages. This huge variety of packages is one of the reasons that R is so successful: the chances are that someone has already solved a problem that you’re working on, and you can benefit from their work by downloading their package.”

From **<http://r-pkgs.had.co.nz/intro.html>** by Hadley Wickham

# Contents of a Package

- R functions
- R data objects
- Help documents for using the package
- Information on the authors, dependencies, etc.
- Information to make sure it “plays well” with R and other packages

# Installing Packages

From CRAN:

```
install.packages("dplyr")
```

From GitHub (for advanced users):

```
library("devtools")
install_github("hadley/dplyr")
```

From Bioconductor (basically CRAN for biology):

```
source("https://bioconductor.org/biocLite.R")
biocLite("qvalue")
```

We will (probably) only be using packages from CRAN. Be *very* careful about dependencies when installing from GitHub.

# Installing Packages (cont'd)

Multiple packages:

```
install.packages(c("dplyr", "ggplot2"))
```

Dependencies:

```
install.packages(c("dplyr", "ggplot2"), dependencies=TRUE)
```

Updating packages:

```
update.packages()
```

# Loading Packages

Two ways to load a package:

```
library("dplyr")
library(dplyr)
```

I prefer the former.

# Getting Started with a Package

When you install a new package and load it, what's next? I like to look at the help files and see what functions and data sets a package has.

```
library("dplyr")
help(package="dplyr")
```

# Specifying a Function within a Package

You can call a function from a specific package. Suppose you are in a setting where you have two packages loaded that have functions with the same name.

```
dplyr::arrange(mtcars, cyl, disp)
```

This calls the `arrange` function specifically from `dplyr`. The package `plyr` also has an `arrange` function.



# More on Packages

We will be covering several highly used R packages in depth this semester, so we will continue to learn about packages, how they are organized, and how they are used.

You can download the “source” of a package from R and take a look at the contents if you want to dig deeper. There are also many good tutorials on creating packages, such as

**<http://hilaryparker.com/2014/04/29/writing-an-r-package-from-scratch/>**.